

# Quiz Review

Peter Kabai

# Things We Will Review

- Basic data science
- Classification, regression and clustering
- KNN
- R coding

# The Basics

- What do we want to make good predictions on?
- Training or test data? Why?
- What happens if our model is too specific to the training data?
- What about if it doesn't fit the training data enough?
- What is a feature? What about a feature space?

# Classification

- We're trying to predict a label, or category.
- This is supervised learning, we know the label of each training example, that is, what category it belongs to.
- Example: given information about credit card debt, we try to predict whether or not the person defaulted on their debt.
- What could some possible target features be categories be?
- If the data shows the output to be 1 or 0, is that correct? Aren't numerical features for regression?
- What is an example of something that cannot be the output of a model that does classification?

# Regression

- Here we're trying to predict a numerical output.
- This is supervised learning, we know the value of the target feature for each training example.
- Example: given the engine size, transmission type, and vehicle weight, predict the mpg.
- What are other possible target features of a model that uses car data?
- What are features that would not be target features in a regression problem?



# Clustering

- Here we're grouping data together into any number of categories.
- Unsupervised learning, we give the model some data, and it finds similarities on its own.

# KNN

- KNN can be used for both regression, and classification.
- In KNN classification the output is the most frequent label of the K closest points.
- In KNN regression the output is the mean of the K closest points.

# KNN Visual





**R Coding!**

# Vectors

```
# There are no scalars! This expression is equal to 1  
length(5)
```

```
## [1] 1
```

```
# The function length() returns the number of items in a vector  
# All item in a vector must be the same type
```

```
# If they aren't the lower types get converted to the higher type  
c(1, 2, "three")
```

```
## [1] "1"      "2"      "three"
```

# Sampling

*# Sampling is a way to create a vector with random values*

```
sample(1:10, 100)
```

```
## Error in sample.int(length(x), size, replace, prob): cannot take a sample larger than the p
```

*# This fails, because there are only 10 options,*

*# but we're trying to pick 100.*

*# To fix it, we can set replace to TRUE*

# Sampling

```
sample(1:10, 100, replace=TRUE)
```

```
## [1] 1 8 9 1 1 4 7 2 2 9 1 4 8 8 5 2 7 10 3 7 2 4 5 2 9 4 2 7 2
## [32] 3 7 4 5 5 9 7 3 3 5 4 3 10 6 6 1 8 3 9 10 4 6 5 1 9 1 8 1 3
## [63] 1 3 2 4 1 9 7 8 5 6 8 8 6 2 5 6 4 8 5 4 10 3 9 10 10 6 5 1 2
## [94] 2 5 3 7 5 6 4
```

*# We can also simply add the TRUE parameter.*

```
sample(1:10, 100, TRUE)
```

```
## [1] 9 4 8 7 10 6 4 10 6 5 9 9 10 1 6 10 5 8 4 6 10 8 6 2 8 3 6 7 7
## [32] 8 2 7 1 1 6 5 1 10 9 6 6 4 9 5 4 3 1 8 7 4 3 1 8 5 4 6 9 3
## [63] 1 3 8 9 3 2 1 4 8 10 6 10 1 8 5 2 3 1 1 6 9 8 5 1 8 8 9 5 4
## [94] 9 10 1 10 3 7 3
```

# Sequences

*# There are multiple ways to do write parameters for sequences.*

```
seq(from=1, to=10, by=2)
```

```
## [1] 1 3 5 7 9
```

```
seq(1, 10, 2)
```

```
## [1] 1 3 5 7 9
```

# Boolean Vectors

```
# Just a vector of true or false values  
# These can be used to index into vectors  
n = c(TRUE, FALSE, TRUE, FALSE)  
m = c(TRUE, FALSE)  
v = c(1:4)
```

# Boolean Vectors

```
# What is the output of both?  
# Which of these uses recycling?  
v[m]
```

```
## [1] 1 3
```

```
v[n]
```

```
## [1] 1 3
```

```
# The first one uses recycling, because length(v) is 4  
# but length(m) is 2
```

# Boolean Vectors

*# Use sequence to create an identical variable*

```
c(1:100)[c(c(TRUE, FALSE, FALSE), c(FALSE, FALSE))]
```

```
## [1] 1 6 11 16 21 26 31 36 41 46 51 56 61 66 71 76 81 86 91 96
```

*# Answer:*

```
seq(1,100,5)
```

```
## [1] 1 6 11 16 21 26 31 36 41 46 51 56 61 66 71 76 81 86 91 96
```

*# A boolean vector is created that then is used to index*

*# the vector c(1:100)*



# Recycling

*# What is the result of this expression?*

```
c(1:2) * c(1:4)
```

```
## [1] 1 4 3 8
```

*# Answer:*

*# The first vector is recycled, so the result is c(1,4,3,8)*

*# What is the result of this expression?*

```
c(1:2) * c(1:5)
```

```
## Warning in c(1:2) * c(1:5): longer object length is not a multiple of shorter object length
```

```
## [1] 1 4 3 8 5
```

*# This runs, but also creates a warning*

# Vector Operations

```
c(1,2,3,4,5,6,7,8,9,10)
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
# Re-write the above vector in a cleaner way
```

```
c(1:10)
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
# What will be the content of y?
```

```
x = c(1:10)
```

```
y = x + 5
```

```
# This is a vectorized operation, so 5 gets added
```

```
# to ALL the values in x, and x becomes c(6,7,8,9,10,11,12,13,14,15)
```

# Vectorized vs. Aggregate Functions

```
x = c(1:10)
```

```
sqrt(x)
```

```
## [1] 1.000000 1.414214 1.732051 2.000000 2.236068 2.449490 2.645751 2.828427 3.000000 3.162278
```

```
mean(x)
```

```
## [1] 5.5
```

```
min(x)
```

```
## [1] 1
```

```
max(x)
```

```
## [1] 10
```

# Vectorized vs. Aggregate Functions

*# What will be the result of this expression?*

```
mean(x) < 5
```

```
## [1] FALSE
```

*# How about this one?*

```
mean(x < 5)
```

```
## [1] 0.4
```

*# What's the result of the intermediate step above?*

```
x < 5
```

```
## [1] TRUE TRUE TRUE TRUE FALSE FALSE FALSE FALSE FALSE FALSE
```

# Vectorized vs. Aggregate Functions

```
# Remember, when aggregator functions are called on boolean  
# vectors, the values are converted to 1 or 0  
v = c(TRUE,TRUE,TRUE,TRUE,FALSE,FALSE,FALSE,FALSE,FALSE,FALSE)  
# What is the output of this expression?  
sum(v)
```

```
## [1] 4
```

```
# And this one?  
sqrt(v)
```

```
## [1] 1 1 1 1 0 0 0 0 0 0
```

# Data Frames (Using Indices)

```
# Using the built in mtcars data frame  
# How do you get the all rows but just columns 3 to 5?  
dat = mtcars[,3:5]  
dat
```

```
##           disp  hp drat  
## Mazda RX4      160.0 110 3.90  
## Mazda RX4 Wag  160.0 110 3.90  
## Datsun 710      108.0  93 3.85  
## Hornet 4 Drive  258.0 110 3.08  
## Hornet Sportabout 360.0 175 3.15  
## Valiant        225.0 105 2.76  
## Duster 360     360.0 245 3.21  
## Merc 240D      146.7  62 3.69  
## Merc 230       140.8  95 3.92  
## Merc 280       167.6 123 3.92  
## Merc 280C      167.6 123 3.92  
## Merc 450SE     275.8 180 3.07  
## Merc 450SL     275.8 180 3.07  
## Merc 450SLC    275.8 180 3.07  
## Cadillac Fleetwood 472.0 205 2.93  
## Lincoln Continental 460.0 215 3.00  
## Chrysler Imperial 440.0 230 3.23
```

# Data Frames (Using Column Names)

*# How about if we want column names?*

```
dat = mtcars[,c("cyl", "hp")]
```

```
dat
```

```
##           cyl  hp
## Mazda RX4      6 110
## Mazda RX4 Wag  6 110
## Datsun 710      4  93
## Hornet 4 Drive  6 110
## Hornet Sportabout 8 175
## Valiant        6 105
## Duster 360      8 245
## Merc 240D       4  62
## Merc 230        4  95
## Merc 280        6 123
## Merc 280C       6 123
## Merc 450SE      8 180
## Merc 450SL      8 180
## Merc 450SLC     8 180
## Cadillac Fleetwood 8 205
## Lincoln Continental 8 215
## Chrysler Imperial 8 230
## Fiat 128        4  66
```

# Data Frames

*# How about the same as before, but only rows 1 to 5?*

```
dat = mtcars[1:5,c("cyl","hp")]
```

```
dat
```

```
##           cyl  hp
## Mazda RX4      6 110
## Mazda RX4 Wag  6 110
## Datsun 710      4  93
## Hornet 4 Drive  6 110
## Hornet Sportabout 8 175
```



# Data Frames

*# Add the mpg column but only show rows where mpg > 20?*

```
dat = mtcars[mtcars$mpg > 20, c("cyl", "hp", "mpg")]
```

```
dat
```

```
##           cyl  hp  mpg
## Mazda RX4      6 110 21.0
## Mazda RX4 Wag  6 110 21.0
## Datsun 710      4  93 22.8
## Hornet 4 Drive  6 110 21.4
## Merc 240D       4  62 24.4
## Merc 230        4  95 22.8
## Fiat 128        4  66 32.4
## Honda Civic     4  52 30.4
## Toyota Corolla  4  65 33.9
## Toyota Corona   4  97 21.5
## Fiat X1-9       4  66 27.3
## Porsche 914-2   4  91 26.0
## Lotus Europa    4 113 30.4
## Volvo 142E      4 109 21.4
```

# Data Frames

*# We're using a boolean vector to index the rows*

```
mtcars$mpg > 20
```

```
## [1] TRUE TRUE TRUE TRUE FALSE FALSE FALSE TRUE TRUE FALSE FALSE FALSE FALSE FALSE FAI
```

```
## [17] FALSE TRUE TRUE TRUE TRUE FALSE FALSE FALSE FALSE TRUE TRUE TRUE FALSE FALSE FAI
```

*# This shows which rows to include, and which rows to ignore*

# CSV to Plot

```
# Read the CSV
```

```
dat_raw = read.csv("data/quizReview.csv")
```

```
dat_raw
```

```
##           Timestamp Wednesday.10.to.11 Wednesday.11.to.12 Wednesday.2.to.3
## 1 2019/02/05 12:08:25 PM PST      I can make it!      I can make it!
## 2 2019/02/05 12:08:27 PM PST                        I can make it!      I can make it!
## 3 2019/02/05 12:17:12 PM PST                                           I can make it!
## 4 2019/02/05 12:22:43 PM PST
## 5 2019/02/05 2:00:43 PM PST
## 6 2019/02/05 2:01:00 PM PST      I can make it!      I can make it!      I can make it!
## 7 2019/02/05 2:01:10 PM PST
## 8 2019/02/05 2:01:19 PM PST                                           I can make it!
## 9 2019/02/05 2:01:29 PM PST                        I can make it!
## 10 2019/02/05 2:01:32 PM PST
## 11 2019/02/05 2:01:35 PM PST      I can make it!
## 12 2019/02/05 2:01:47 PM PST      I can make it!      I can make it!
## 13 2019/02/05 2:01:57 PM PST
## 14 2019/02/05 2:01:58 PM PST                        I can make it!      I can make it!
## 15 2019/02/05 2:02:02 PM PST      I can make it!      I can make it!
## 16 2019/02/05 2:02:08 PM PST
## 17 2019/02/05 2:02:24 PM PST                                           I can make it!
## 18 2019/02/05 2:02:46 PM PST
```

# CSV to Plot

```
# Remove the first row, which is just a timestamp
```

```
dat = dat_raw[,-1]
```

```
# Convert all values to either a 1 or a 0
```

```
dat[,] = ifelse (is.na(dat) | dat != "I can make it!", 0, 1)
```

```
dat
```

```
##      Wednesday.10.to.11 Wednesday.11.to.12 Wednesday.2.to.3 Wednesday.3.to.4 Thursday.10.to.11
## 1              1              1              0              0              (
## 2              0              1              1              0              (
## 3              0              0              1              1              (
## 4              0              0              0              0              1
## 5              0              0              0              0              (
## 6              1              1              1              1              (
## 7              0              0              0              0              (
## 8              0              0              1              1              (
## 9              0              1              0              0              1
## 10             0              0              0              0              (
## 11             1              0              0              0              1
## 12             1              1              0              0              (
## 13             0              0              0              0              (
## 14             0              1              1              1              (
```

# CSV to Plot

```
# Check to see if all columns are now numeric
```

```
sapply(dat, class)
```

```
## Wednesday.10.to.11 Wednesday.11.to.12 Wednesday.2.to.3 Wednesday.3.to.4 Thursday.10.to.
##           "numeric"           "numeric"           "numeric"           "numeric"           "numeric"
## Thursday.11.to.12 Thursday.12.to.1 Thursday.1.to.2
##           "numeric"           "numeric"           "numeric"
```

```
# Change the names of the columns (gsub wasn't covered in class)
```

```
names(dat) = gsub("[.]", " ", names(dat))
```

```
# Show the number of "I can make it!" votes for each day
```

```
sapply(dat, sum)
```

```
## Wednesday 10 to 11 Wednesday 11 to 12 Wednesday 2 to 3 Wednesday 3 to 4 Thursday 10 to
##           5           8           7           7
## Thursday 11 to 12 Thursday 12 to 1 Thursday 1 to 2
##           6           10           13
```

# CSV to Plot

```
# Sorting that vector
```

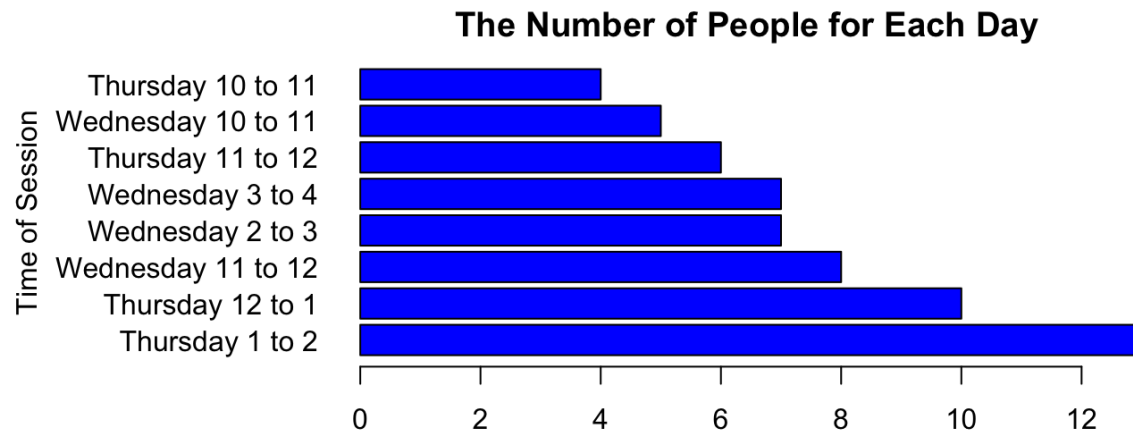
```
vect_of_times = sort(sapply(dat, sum), decreasing = TRUE)
```

```
vect_of_times
```

```
##      Thursday 1 to 2      Thursday 12 to 1 Wednesday 11 to 12 Wednesday 2 to 3 Wednesday 3 to 4
##              13              10              8              7
## Thursday 11 to 12 Wednesday 10 to 11 Thursday 10 to 11
##              6              5              4
```

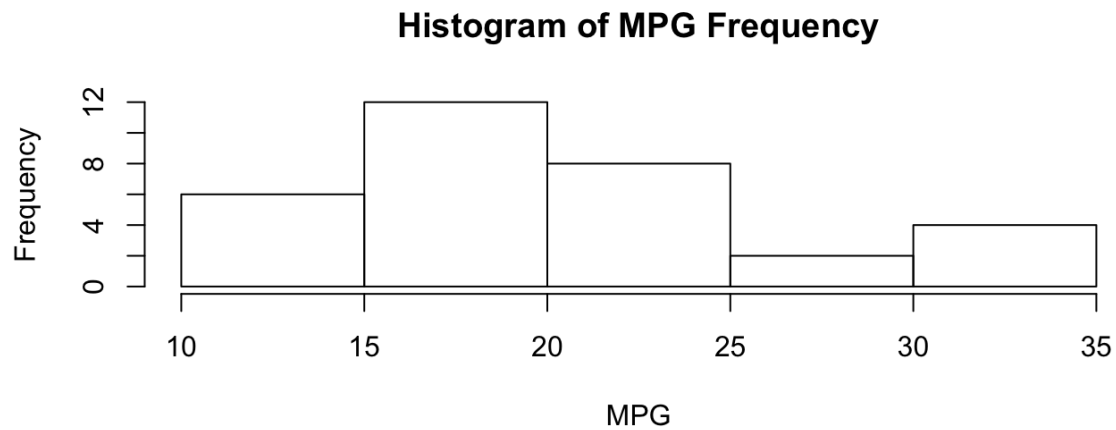
# CSV to Plot

```
# plot the named vector of times
par(mar = c(4, 11, 2, 1), mgp = c(9,1,0))
barplot(
  vect_of_times, las = 1, horiz = TRUE, col = "blue",
  ylab = "Time of Session", main = "The Number of People for Each Day"
)
```



# Other Types of Plots (Histogram)


```
hist(  
  mtcars$mpg,  
  main = "Histogram of MPG Frequency", xlab = "MPG"  
)
```





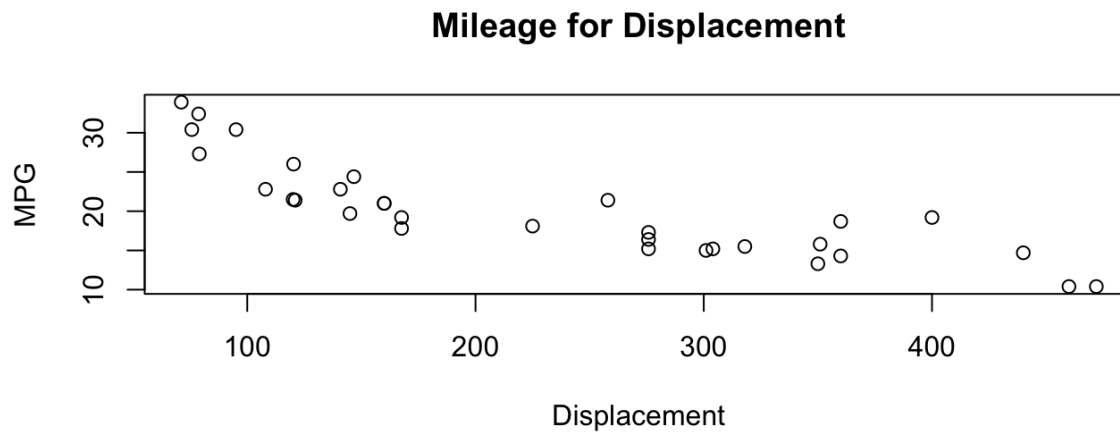
# Other Types of Plots (Line)

```
x = sample(1:10, 10, TRUE) # vector of random values
plot(
  x, type="b",
  main = "Random Numbers in a Vector", xlab = "Index", ylab = "Random Value"
)
```



# Other Types of Plots (Scatter)

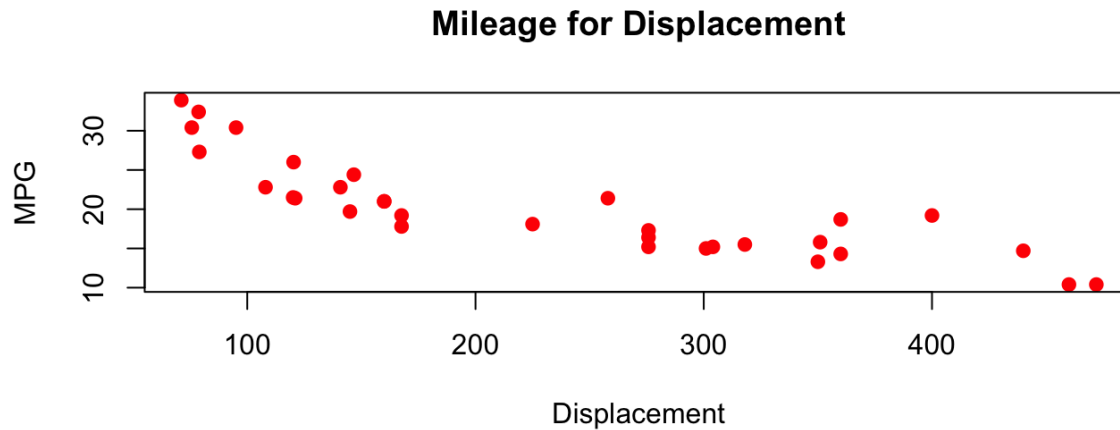
```
plot(  
  mtcars$disp, mtcars$mpg,  
  main = "Mileage for Displacement", xlab = "Displacement", ylab = "MPG"  
)
```



# Other Types of Plots (Scatter)

*# a different way to do the same thing*

```
plot(  
  mpg ~ disp, data=mtcars, col="red", pch=19,  
  main = "Mileage for Displacement", xlab = "Displacement", ylab = "MPG"  
)
```



Anything else?